

Concurrent Programming Notes

v0.015 – January 19, 2023

Eduardo Bonelli

Preface

These course notes provide supporting material for CS511.

Contents

Preface	i
1 Shared Memory Model and Transition Systems	1
1.1 Shared Memory Model	1
1.2 Transition Systems	3
1.3 Atomicity	6
1.4 The Mutual Exclusion Problem	8
2 Semaphores	9
2.1 Introduction	9
2.2 The MEP Problem Revisited	9
2.3 More Examples	11
2.3.1 Thread Dumps	12
2.4 Classical Synchronization Problems	14
2.4.1 Producers/Consumers	14
2.4.2 Readers/Writers	16
2.4.3 Barrier Synchronization	16
3 Monitors	21
3.1 A monitor implementing a semaphore	22
3.2 Producers/Consumers	23
3.3 Readers/Writers	24
4 Promela	29
4.1 Syntax	29
4.1.1 Examples involving Loops	30
4.1.2 Expressions as blocking commands	30
4.1.3 Macros	32
4.2 Assertion-Based Model Checking	35
4.2.1 The Bar Problem Revisited	35
4.2.2 The MEP Problem	38
4.3 Non-Progress Cycles	40
4.3.1 The Zoo Problem Revisited	44
5 Solution to Selected Exercises	47

Chapter 1

Shared Memory Model and Transition Systems

This chapter...

1.1 Shared Memory Model

We begin with an example of a program in Groovy.

```
1  int x = 0
3  Thread.start { //P
    x = 1
5  }
7  Thread.start { //Q
    x = 2
9  }
```

ex1.groovy

This program declares a shared variable `x`, sets it to 0 and then spawns two threads. The first thread sets `x` to 1 and the second to 2. After this program terminates, the value of `x` may either be 1 or 2. The variable `x` is said to be shared in the sense that it is visible to (or its scope includes) both threads¹.



Semicolons are optional in Groovy

Assuming this program is stored in a file called `ex1.groovy`, it may be executed using the terminal as follows:

¹From the point of view of Groovy it is actually a local variable. In Groovy, global variables are declared by omitting the type annotation.

```
$ groovy ex1
2 $
```

bash

Since our program contains no output statements, there is no visible effect from its execution. The following example, waits for P and Q to terminate using the built-in method `join` and then prints the value of `x`:

```
int x = 0
2
P = Thread.start { //P
4     x = 1
}
6
Q = Thread.start { //Q
8     x = 2
}
10
P.join() // Wait for P to terminate
12 Q.join() // Wait for Q to terminate
println x
```

ex2.groovy

Assuming this program is stored in a file called `ex2.groovy`, it may be executed using the terminal as follows:

```
$ groovy ex2
2 2
$
```

bash

Repeated execution will most likely produce 2 since P is spawned before Q and runs immediately. It is entirely possible, however, to obtain 1 as a result.

The following example is a Groovy program that prints characters.

```
1 Thread.start { //P
    print "A"
3     print "B"
}
5
Thread.start { //Q
7     print "C"
}
```

What are the possible outputs one may obtain from executing it? It can print three possible sequences of characters, namely ABC, ACB, CAB. What about the following program?

```
2 Thread.start { //P
    print "A"
    print "B"
4 }
6 Thread.start { //Q
    print "C"
8     print "D"
```



```
}

```

Clearly the number of possible executions, also called interleavings, grows exponentially with the number of instructions in each thread. Indeed, if P has m instructions and Q has n instructions, then there are

$$\binom{m+n}{m} = \frac{(m+n)!}{m!n!}$$

This makes it difficult to reason about concurrent programs: there are simply too many interleavings to consider; we never know whether one such interleaving might lead our code to produce an unwanted result. We clearly need some rigorous device to be able to model all such possible interleavings and check whether they satisfy our intended properties. A device that describes the run-time execution of a concurrent program. There is a further, equally important reason, why we need this device. Consider the following program:

```
1 int x=0 // shared variable
3 P = Thread.start {
    x = x+1
5 }
Q = Thread.start {
    x = x+1
7 }
9
P.join() // wait for P to terminate
11 Q.join() // wait for Q to terminate
13 println(x)
```

Its execution produces 1 as output!

```
1 $ groovy ex1
1
3 $
```

bash

How is that possible?

1.2 Transition Systems

This section introduces transition systems, a device we use to model the run-time behavior of concurrent programs. After defining transition systems, we illustrate how to associate a transition system to Groovy programs. By doing so, we assign “meaning” to our concurrent programs. It should be mentioned that we will associate transition systems only to a subset of simple Groovy programs, not arbitrary ones.

A **Transition System** \mathcal{A} is a tuple (S, \rightarrow, I) where

- S is a set of states;
- $\rightarrow \subseteq S \times S$ is a transition relation; and
- $I \subseteq S$ is a set of initial states.

We say that \mathcal{A} is finite if S is finite. Also, we write $s \rightarrow s'$ for $(s, s') \in \rightarrow$.


We illustrate, in this first example, how to model the runtime execution of Example 1.1, repeated below:

```

1  int x = 0
3  Thread.start { //P
    x = 1
5  }
7  Thread.start { //Q
    x = 2
9  }

```

The states of our transition system will consist of 3-tuples containing the instruction pointer for p , the instruction pointer for q and the value of x . The initial state is signalled with a small arrow. The hyphen indicates that there are no further instructions to be executed by that thread.

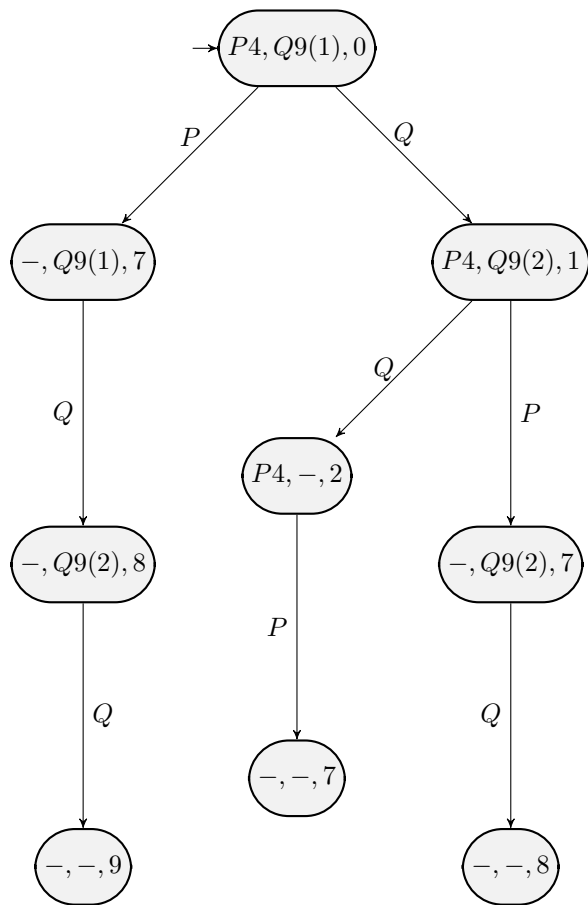


How we hardcode “for”-loops with a constant upper bound.

```

1  int x = 0
2
3  Thread.start { //P
4      x = 7
5  }
6
7  Thread.start { //Q
8      2.times {
9          x = x+1
10     }
11 }

```



How we distinguish local variables with the same name using "local_P" and "local_Q" in the state format.

```

1  int x = 0 // shared variable
3  Thread.start { //P
    int local = x
5   x = local+1 // atomic
  }
7
  Thread.start { //Q
9   int local = x
    x = local+1 // atomic
11 }

```

How we deal with "while"-loops.

```

1  int x=0 // shared variable
3  Thread.start { //P
    while (x<1) {
5     print x
    }

```

```

7 }
9 Thread.start { //Q
    x = x + 1
11 }

```

How we decorate transitions with the output string of a print

```

1 int x = 0 // shared variable
3 Thread.start { //P
    x = x + 1
5     x = x + 1
    }
7
9 Thread.start { //Q
    while (x >= 2)
        print x
11 }

```

```

1 int counter=0 // shared variable
3 P = Thread.start {
    50.times {
5         counter = counter+1
        }
7     }
9 Q = Thread.start {
    50.times {
        counter= counter+1
11     }
    }
13
15 P.join() // wait for P to finish
    Q.join() // wait for Q to finish
17
19 println counter // print value of counter

```

1.3 Atomicity

Consider the following program:

```

1 x=0
3 Thread.start { //P
    x = x + 1
5     println x
    }
7
9 Thread.start { //Q
    x = x + 1
11     println x
    }

```

One would expect 1 and 2, or 2 and 2 to be printed. These are indeed possible outputs. However, 1 and 1 is also possible:

```

1 $ groovy ex3
1
3 1
$

```

bash

The reason is that assignment is not an atomic operation, rather it is decomposed into more fine grained (bytecode) operations. It is the latter that are interleaved. Let's take a closer look at those fine grained operations. Consider the following Java class that spawn two threads, each of which updates a shared variable:

```

class A implements Runnable {
2
    static int x=0;
4
    public void run() {
6
        x=x+1;
    }
8
    public static void main(String[] args) {
10
        new Thread(new A()).start();
        new Thread(new A()).start();
12
    }
}

```

A.java

We compile it and look at the resulting bytecode by using `javap`, the Java class file disassembler:

```

$ javac A.java
$ javap -c A
Compiled from "A.java"
4 class A implements java.lang.Runnable {
    static int x;
6
    A();
8    Code:
        0: aload_0
        1: invokespecial #1                  // Method java/lang/Object."<init>":()V
        4: return
12
    public void run();
14    Code:
        0: getstatic     #7                  // Field x:I
        3: iconst_1
        4: iadd
        5: putstatic     #7                  // Field x:I
        8: return
16
    public static void main(java.lang.String[]);
20    Code:
        0: new           #13                 // class java/lang/Thread
        3: dup
22

```

```

26      4: new          #8          // class A
      7: dup
      8: invokespecial #15          // Method "<init>":()V
28     11: invokespecial #16          // Method java/lang/Thread."<init>":(Ljava
      14: invokevirtual #19          // Method java/lang/Thread.start:()V
30     17: new          #13          // class java/lang/Thread
      20: dup
32     21: new          #8          // class A
      24: dup
34     25: invokespecial #15          // Method "<init>":()V
      28: invokespecial #16          // Method java/lang/Thread."<init>":(Ljava
36     31: invokevirtual #19          // Method java/lang/Thread.start:()V
      34: return
38
40     static {};
      Code:
      0: iconst_0
42     1: putstatic      #7          // Field x:I
      4: return
44 }

```

bash

The only lines we are interested are lines 15 to 18. Each thread has a JVM stack. Every time a method is called, a new frame is created (heap-allocated) and stored on the JVM stack for that thread. Each frame has its own array of local variables, its own operand stack, and a reference to the run-time constant pool of the class of the current method. The instruction $x=x+1$ is compiled to four bytecode instructions whose meaning can be read off from their opcodes:

```

2      0: getstatic      #7          // Field x:I
      3: iconst_1
      4: iadd
4      5: putstatic      #7          // Field x:I

```

It is these operations, for each thread, that get interleaved. Thus, it is possible to have the following interleaving:

```

2      0(P): getstatic #7          // Field x:I
      0(Q): getstatic #7          // Field x:I
      3(P): iconst_1
4      3(Q): iconst_1
      4(P): iadd
6      4(Q): iadd
      5(P): putstatic #7          // Field x:I
8      5(Q): putstatic #7          // Field x:I

```

These instructions end up storing 1 in x .

1.4 The Mutual Exclusion Problem

Chapter 2

Semaphores

2.1 Introduction

2.2 The MEP Problem Revisited

Consider the following solution to the MEP problem using a binary semaphore presented in listing 2.1¹.

```
1 Semaphore mutex= new Semaphore(1)
2
3 Thread.start { //P
4     while (true) {
5         mutex.acquire()
6         mutex.release()
7     }
8 }
9 Thread.start { //Q
10    while (true) {
11        mutex.acquire()
12        mutex.release()
13    }
14 }
```

Listing 2.1: Solution to MEP using a binary semaphore

One easy way to verify that all three properties of MEP are upheld is to construct its transition system and then analyze these properties. This requires a means for representing semaphores. Since a semaphore is an object with state and the latter includes the number of permits and the set of blocked processes, we shall model `mutex` using the expression `mutex[i,s]` where `i` is the number of permits and `s` is a set of blocked processes. Moreover, we use the “!” symbol as instruction pointer in the states of our transition systems to indicate that there are no instructions ready to execute. For example, a state such as $P6,!,mutex[0,\{Q11\}]$, reflects that only P can be scheduled for execution, there are no permits available in `mutex` and one thread is blocked on

¹Groovy requires that you import the Semaphore class in order to be able to use it. All code excerpts involving semaphores should thus include, at the top, the line `import java.util.concurrent.Semaphore`. This is typically omitted in our examples.




Figure 2.1: Transition System for the solution to MEP using a binary semaphore

`mutex` waiting for a permit to become available, namely `Q`. Figure ?? is the transition system for the listing in Figure 2.1.

Consider the setting where the above solution is applied to three threads wanting to access their CS. This is illustrated in Listing 2.2. Although `Mutex` and `Absence of Livelock` are upheld, `Freedom From Starvation` is not. Indeed, consider the scenario where `P` goes in and `Q` and `R` try to get in and are both blocked and placed in the set of blocked processes for `mutex`. [COMPLETE]??

This is easily solved by having the set of blocked processes in `mutex` be a queue. Such semaphores are called fair semaphores. This is achieved by using an alternative constructor for semaphores that includes a fairness parameter

```
Semaphore(int permits, boolean fair)
```

Replacing line 1 in Listing 2.2 with `Semaphore mutex= new Semaphore(1,true)` suffices to obtain a correct solution to the MEP for any number of threads.

```

1 Semaphore mutex= new Semaphore(1)
2
3 Thread.start { //P
4     while (true) {
5         mutex.acquire()
6         mutex.release()
7     }
8 }
9 Thread.start { //Q
10    while (true) {
11        mutex.acquire()
12        mutex.release()
13    }

```



```

14 }
15 Thread.start { //R
16     while (true) {
17         mutex.acquire()
18         mutex.release()
19     }
20 }

```

Listing 2.2: Attempt at solving the MEP using a binary semaphore for $N=3$

2.3 More Examples

```

Thread.start {
2     println "A"
    println "B"
4 }
Thread.start {
6     println "C"
    println "D"
8 }

```

```

Semaphore cAfterA = new Semaphore(0)
2
Thread.start {
4     println "A"
    mutex.release()
6     println "B"
}
8 Thread.start {
    mutex.acquire()
10    println "C"
    println "D"
12 }

```

Consider the following example which prints any (infinite) sequence of “a”s and “b”s:

```

Thread.start { //P
2     while (true) {
        print "a"
4     }
}
6
Thread.start { //Q
8     while (true) {
        print "b"
10    }
}

```

Using semaphores, how would you ensure that only the infinite sequence “aabaabaab...” is printed? Hint: make use of two semaphores, `a` and `b`, enabling the execution of an iteration in `P` and an iteration in `Q`, respectively.

Here is a solution.

```

import java.util.concurrent.Semaphore
2

```

```

Semaphore a = new Semaphore(2)
4 Semaphore b = new Semaphore(0)

6 Thread.start { //P
    while (true) {
8         a.acquire()
        print "a"
10        b.release()
    }
12 }

14 Thread.start { //Q
    while (true) {
16        b.acquire(2)
        print "b"
18        a.release(2)
    }
20 }

```

2.3.1 Thread Dumps

We can check the current thread dump of the our Groovy/Java application as follows. Let's use the example above. First we modify our code so that we give our threads an easy to spot name and remove the lines that print. The result is below; we'll call it `ex1.groovy`.

```

import java.util.concurrent.Semaphore

2 Semaphore a = new Semaphore(2)
3 Semaphore b = new Semaphore(0)

6 Thread.start { //P
    Thread.currentThread().setName("P Thread");
8    while (true) {
        a.acquire()
10        // print "a"
        b.release()
    }
12 }

14 Thread.start { //Q
16    Thread.currentThread().setName("Q Thread");
    while (true) {
18        b.acquire()
        b.acquire()
20        // print "b"
        a.release()
22        a.release()
    }
24 }

```

Now we run it in the background, use `jstack` to obtain the stack trace of the bash job and send the output to a text file `thead-dump.txt`

```

$ groovy ex1 &
2 [1] 23275
$ jstack -l 23275 > thread-dump.txt

```

```

4 $ kill %1
    [1] + 23275 exit 143  groovy ex1
6 $ emacs thread-dump.txt

```

bash

The dump contains information on all threads involved in our application. We'll just show an excerpt that mentions `p` and `q`. We can see that the former is in a `RUNNABLE` state and the latter is in a `WAITING` state. We can also see the current instruction being executed in each thread.

```

"P Thread" #17 prio=5 os_prio=31 cpu=5910.73ms elapsed=10.91s tid=0x00007f7d9412b400 nid=27655 runnable
2  java.lang.Thread.State: RUNNABLE
    at jdk.internal.misc.Unsafe.unpark(java.base@18.0.1.1/Native Method)
4  at java.util.concurrent.locks.LockSupport.unpark(java.base@18.0.1.1/LockSupport.java:177)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.signalNext(java.base@18.0.1.1/AbstractQueuedSynchronizer.java:432)
6  at java.util.concurrent.locks.AbstractQueuedSynchronizer.releaseShared(java.base@18.0.1.1/AbstractQueuedSynchronizer.java:432)
    at java.util.concurrent.Semaphore.release(java.base@18.0.1.1/Semaphore.java:432)
8  at java.lang.invoke.LambdaForm$DMH/0x0000000800d28000.invokeVirtual(java.base@18.0.1.1/LambdaForm$DMH/0x0000000800d28000)
    at java.lang.invoke.LambdaForm$MH/0x0000000800e32c00.invoke(java.base@18.0.1.1/LambdaForm$MH/0x0000000800e32c00)
10 at java.lang.invoke.LambdaForm$MH/0x0000000800e2b400.guardWithCatch(java.base@18.0.1.1/LambdaForm$MH/0x0000000800e2b400)
    at java.lang.invoke.DelegatingMethodHandle$Holder.delegate(java.base@18.0.1.1/DelegatingMethodHandle$Holder.java:318)
12 at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH/0x0000000800e27800)
    at java.lang.invoke.DelegatingMethodHandle$Holder.delegate(java.base@18.0.1.1/DelegatingMethodHandle$Holder.java:318)
14 at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH/0x0000000800e27800)
    at java.lang.invoke.Invokers$Holder.linkToCallSite(java.base@18.0.1.1/Invokers$Holder.java:258)
16 at ex1$_run_closure1.doCall(ex1.groovy:12)
    at ex1$_run_closure1.doCall(ex1.groovy)
18 at java.lang.invoke.DirectMethodHandle$Holder.invokeSpecial(java.base@18.0.1.1/DirectMethodHandle$Holder.java:318)
    at java.lang.invoke.LambdaForm$MH/0x0000000800c1c800.invoke(java.base@18.0.1.1/LambdaForm$MH/0x0000000800c1c800)
20 at java.lang.invoke.Invokers$Holder.invokeExact_MT(java.base@18.0.1.1/Invokers$Holder.java:258)
    at jdk.internal.reflect.DirectMethodHandleAccessor.invokeImpl(java.base@18.0.1.1/DirectMethodHandleAccessor.java:117)
22 at jdk.internal.reflect.DirectMethodHandleAccessor.invoke(java.base@18.0.1.1/DirectMethodHandleAccessor.java:117)
    at java.lang.reflect.Method.invoke(java.base@18.0.1.1/Method.java:577)
24 at org.codehaus.groovy.reflection.CachedMethod.invoke(CachedMethod.java:343)
    at groovy.lang.MetaMethod.doMethodInvoke(MetaMethod.java:328)
26 at org.codehaus.groovy.runtime.metaclass.ClosureMetaClass.invokeMethod(ClosureMetaClass.java:279)
    at groovy.lang.MetaClassImpl.invokeMethod(MetaClassImpl.java:1009)
28 at groovy.lang.Closure.call(Closure.java:418)
    at groovy.lang.Closure.call(Closure.java:412)
30 at groovy.lang.Closure.run(Closure.java:500)
    at java.lang.Thread.run(java.base@18.0.1.1/Thread.java:833)
32
    Locked ownable synchronizers:
34      - None

"Q Thread" #18 prio=5 os_prio=31 cpu=6110.53ms elapsed=10.91s tid=0x00007f7d9411fa00 nid=28163 runnable
36  java.lang.Thread.State: WAITING (parking)
    at jdk.internal.misc.Unsafe.park(java.base@18.0.1.1/Native Method)
38      - parking to wait for <0x00000006180b9960> (a java.util.concurrent.Semaphore$NonfairSync)
40  at java.util.concurrent.locks.LockSupport.park(java.base@18.0.1.1/LockSupport.java:211)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.acquire(java.base@18.0.1.1/AbstractQueuedSynchronizer.java:432)
42  at java.util.concurrent.locks.AbstractQueuedSynchronizer.acquireSharedInterruptibly(java.base@18.0.1.1/AbstractQueuedSynchronizer.java:432)
    at java.util.concurrent.Semaphore.acquire(java.base@18.0.1.1/Semaphore.java:318)
44  at java.lang.invoke.LambdaForm$DMH/0x0000000800d28000.invokeVirtual(java.base@18.0.1.1/LambdaForm$DMH/0x0000000800d28000)
    at java.lang.invoke.LambdaForm$MH/0x0000000800e32c00.invoke(java.base@18.0.1.1/LambdaForm$MH/0x0000000800e32c00)
46  at java.lang.invoke.LambdaForm$MH/0x0000000800e2b400.guardWithCatch(java.base@18.0.1.1/LambdaForm$MH/0x0000000800e2b400)
    at java.lang.invoke.DelegatingMethodHandle$Holder.delegate(java.base@18.0.1.1/DelegatingMethodHandle$Holder.java:318)

```

```

48  at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH
   at java.lang.invoke.DelegatingMethodHandle$Holder.delegate(java.base@18.0.1.1/DelegatingMet
50  at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH
   at java.lang.invoke.Invokers$Holder.linkToCallSite(java.base@18.0.1.1/Invokers$Holder)
52  at ex1$_run_closure2.doCall(ex1.groovy:19)
   at ex1$_run_closure2.doCall(ex1.groovy)
54  at java.lang.invoke.DirectMethodHandle$Holder.invokeSpecial(java.base@18.0.1.1/DirectMethod
   at java.lang.invoke.LambdaForm$MH/0x0000000800c1c800.invoke(java.base@18.0.1.1/LambdaForm$M
56  at java.lang.invoke.Invokers$Holder.invokeExact_MT(java.base@18.0.1.1/Invokers$Holder)
   at jdk.internal.reflect.DirectMethodHandleAccessor.invokeImpl(java.base@18.0.1.1/DirectMeth
58  at jdk.internal.reflect.DirectMethodHandleAccessor.invoke(java.base@18.0.1.1/DirectMethodHa
   at java.lang.reflect.Method.invoke(java.base@18.0.1.1/Method.java:577)
60  at org.codehaus.groovy.reflection.CachedMethod.invoke(CachedMethod.java:343)
   at groovy.lang.MetaMethod.doMethodInvoke(MetaMethod.java:328)
62  at org.codehaus.groovy.runtime.metaclass.ClosureMetaClass.invokeMethod(ClosureMetaClass.jav
   at groovy.lang.MetaClassImpl.invokeMethod(MetaClassImpl.java:1009)
64  at groovy.lang.Closure.call(Closure.java:418)
   at groovy.lang.Closure.call(Closure.java:412)
66  at groovy.lang.Closure.run(Closure.java:500)
   at java.lang.Thread.run(java.base@18.0.1.1/Thread.java:833)
68
   Locked ownable synchronizers:
70   - None

```

One could also make use of online tools that analyse these thread dumps to help identify potential issues. For example, you can try and upload `thread-dump.txt` to this site fastthread.io and click on “analyze”.

2.4 Classical Synchronization Problems

This section addresses some classical synchronization problems using semaphores.

2.4.1 Producers/Consumers

Buffer of size 1, one producer and one consumer. The code below also works if there were multiple producers and multiple consumers.

```

Integer buffer
2
Semaphore consume = new Semaphore(0)
3 Semaphore produce = new Semaphore(1)
4
6 Thread.start { // Prod
   Random r = new Random()
7   while (true) {
8     produce.acquire()
9     buffer = r.nextInt(10000) // produce()
10    println "produced "+buffer
11    Thread.sleep(1000)
12    consume.release()
13  }
14 }
15
16 Thread.start { // Cons
17

```

```

20     while (true) {
        consume.acquire()
        println "consumed "+buffer
22     buffer = null // consume(buffer)
        produce.release()
24     }
    }

```

Buffer of size N with one producer and one consumer. Also known as a blocking queue.

```

final int N=10
2 Integer[] buffer = [0] * N

4 Semaphore consume = new Semaphore(0)
Semaphore produce = new Semaphore(N)
6 int start = 0
int end = 0

8 Thread.start { // Prod
10     Random r = new Random()
    while (true) {
12         produce.acquire()
        mutexP.acquire()
14         buffer[start] = r.nextInt(10000) // produce()
        println id+" produced "+buffer[start] + " at index "+start
16         start = (start + 1) % N
        mutexP.release()
18         consume.release()
    }
20 }

22 Thread.start { // Cons
    while (true) {
24         consume.acquire()
        mutexC.acquire()
26         println id+ " consumed "+buffer[end] + " at index "+end
        buffer[end] = null // consume(buffer)
28         end = (end + 1) % N
        mutexC.release()
30         produce.release()
    }
32 }

```

Buffer of size N with multiple producers and multiple consumers.



The static method `currentMethod()` returns a reference to the currently executing thread object. Every thread has a unique id. It may be obtained by using the `getId()` method.

```

final int N=10
2 Integer[] buffer = [0] * N

4 Semaphore consume = new Semaphore(0)
Semaphore produce = new Semaphore(N)
6 Semaphore mutexP = new Semaphore(1) // mutex to avoid race conditions on start
Semaphore mutexC = new Semaphore(1) // mutex to avoid race conditions on end

```

```

8  int start = 0
   int end = 0
10
12  5.times {
13      Thread.start { // Prod
14          Random r = new Random()
15          while (true) {
16              produce.acquire()
17              mutexP.acquire()
18              buffer[start] = r.nextInt(10000) // produce()
19              println Thread.currentThread().getId()+" produced "+buffer[start] + " at index "+start
20              start = (start + 1) % N
21              mutexP.release()
22              consume.release()
23          }
24      }
25  }
26
27  5.times{
28      Thread.start { // Cons
29          while (true) {
30              consume.acquire()
31              mutexC.acquire()
32              println Thread.currentThread().getId()+" consumed "+buffer[end] + " at index "+end
33              buffer[end] = null // consume(buffer)
34              end = (end + 1) % N
35              mutexC.release()
36              produce.release()
37          }
38      }
39  }

```

2.4.2 Readers/Writers

2.4.3 Barrier Synchronization

```

// One-time use barrier
// Barrier size = N
// Total number of threads in the system = N
4
final int N=3
6  N.times {
7      Thread.start {
8          while (true) {
9              // barrier arrival protocol
10
11              // barrier
12          }
13      }
14  }

```

```

import java.util.concurrent.Semaphore
2  // One-time use barrier
// Barrier size = N
4  // Total number of threads in the system = N

```

```

6  final int N=3
   int t=0
8  Semaphore barrier = new Semaphore(0)
   Semaphore mutex = new Semaphore(1)
10 N.times {
    Thread.start {
12         while (true) {
            // barrier arrival protocol
14             mutex.acquire()
            if (t<N) {
16                 t++
                if (t==N) {
18                     N.times { barrier.release() }
                }
20             } else {
                barrier.release()
22             }
            mutex.release()
24             // barrier
            barrier.acquire()
26         }
    }
28 }

```

Using cascaded signalling:

```

import java.util.concurrent.Semaphore
2 // One-time use barrier
   // Barrier size = N
4 // Total number of threads in the system = N

6 final int N=3
   int t=0
8 Semaphore barrier = new Semaphore(0)
   Semaphore mutex = new Semaphore(1)
10 N.times {
    Thread.start {
12         while (true) {
            // barrier arrival protocol
14             mutex.acquire()
            if (t<N) {
16                 t++
                if (t==N) {
18                     barrier.release()
                }
20             }
            mutex.release()
22             // barrier
            barrier.acquire() // Cascaded signalling
24             barrier.release()
        }
    }
26 }

```

Cyclic (or reusable) barrier. Failed attempt:

```

1 import java.util.concurrent.Semaphore

```

```

3 // Cyclic (ie. Reusable) barrier
// Barrier size = N
// Total number of threads in the system = N
5
Semaphore mutex = new Semaphore(1)
7 Semaphore barrier = new Semaphore(0)
final int N = 3
9 int t=0
11 N.times {
    Thread.start {
13         while (true) {
            // arrival
15             mutex.acquire()
            t++;
17             if (t==N) {
                N.times { barrier.release()}
19                 t=0 // attempt to reset barrier counter
            }
21             mutex.release()
23             // barrier
            barrier.acquire()
25         }
    }
27 }

```

One easy way to verify that it is incorrect is to count the number of times a thread cycles passed the barrier. Then, notice that some threads can race far ahead of others in terms of the difference in number cycles; this difference can be larger than 1.

A solution follows. We use a second barrier to wait for all threads to fall through the first barrier, thus avoiding any one thread getting ahead of the others.

```

1 import java.util.concurrent.Semaphore
3 // Cyclic (ie. Reusable) barrier
// Barrier size = N
5 // Total number of threads in the system = N
7 Semaphore mutex = new Semaphore(1)
Semaphore barrier = new Semaphore(0)
9 Semaphore barrier2 = new Semaphore(0)
final int N = 3
11 int t=0
13 N.times {
    int id = it
15    Thread.start {
        while (true) {
17            // arrival
            mutex.acquire()
19            t++;
            if (t==N) {
21                N.times { barrier.release() }
            }
23            mutex.release()

```



```
25     // barrier
26     barrier.acquire()
27
28     mutex.acquire()
29     t--
30     if (t==0) {
31         N.times { barrier2.release() }
32     }
33     mutex.release()
34
35     barrier2.acquire()
36 }
37
38 }
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
```

43 }

Chapter 3

Monitors

A monitor is a program module that encapsulates data and operations and, moreover, guarantees mutual exclusion in the execution of the operations.

Listing 3.1 implements two turnstiles each of which accesses a global counter. The counter is implemented using a monitor. This monitor supports operations `inc()`, `dec()` and `read()`. The `synchronized` qualifier ensures mutual exclusion in the execution of these methods. Every object has a built in lock called an intrinsic lock. When a thread invokes a synchronized method, it automatically acquires the intrinsic lock for that method's object and releases it when the method returns. A thread is said to own the intrinsic lock between the time it has acquired the lock and released the lock. As long as a thread owns an intrinsic lock, no other thread can acquire the same lock. The other thread will block when it attempts to acquire the lock.

```
// Monitor declaration
2  class Counter {
    private int c
4
    public synchronized void inc() {
6        c++
    }
8
    public synchronized void dec() {
10        c--
    }
12
    public synchronized int read() {
14        return c
    }
16 }

18 // Sample use of the monitor
Counter ctr = new Counter()
20
22 P = Thread.start {
    10.times {
24         ctr.inc()
    }
26 }
```

```

Q = Thread.start {
28     10.times {
        ctr.inc()
30     }
    }
32
P.join()
34 Q.join()
println (ctr.read())

```

Listing 3.1: Avoiding race conditions on a shared counter using a monitor

3.1 A monitor implementing a semaphore

```

1  class Semaphore {
    private int permits
3
    Semaphore(int init) {
5        permits=init
    }
7
    public synchronized void acquire() {
9        while (permits==0) {
            wait()
11       }
        permits--
13    }
15
    public synchronized void release() {
        notify()
17        permits++
    }
19 }
21
Semaphore mutex = new Semaphore(1)
int c=0
23
P = Thread.start {
25     10.times {
        mutex.acquire()
27         c++
        mutex.release()
29     }
    }
31
Q = Thread.start {
33     10.times {
        mutex.acquire()
35         c++
        mutex.release()
37     }
    }
39
P.join()
41 Q.join()
println c

```

This solution is not fair on the threads that are sleeping since an outside thread could “steal” the permit what is made available through a call to `release`. An alternative that is fair in this sense¹ is given in Listing 3.2. A different approach is followed in [Car96].

```

import java.util.concurrent.Semaphore

2
class Semaphore {
4     private int permits;
     private long startWaitingTime=0;
6     private static final long startTime = System.currentTimeMillis();
     private int waiting=0;

8
     Semaphore(int n) {
10         permits=n;
     }

12
     synchronized protected static final long age() {
14         return System.currentTimeMillis() - startTime;
     }

16
     synchronized void acquire() {
18         if (waiting>0 || permits==0) {
             long arrivalTime = age();
20             while (arrivalTime>startWaitingTime || permits==0) {
                 waiting++;
22                 wait();
                 waiting--;
24             }
         }
26         permits--;
     }

28
     synchronized void release() {
30         permits++;
         startWaitingTime = age();
32         notify();
     }
34 }

```

Listing 3.2: Fair semaphores

3.2 Producers/Consumers

```

class PC {
2     private Object buffer;

4     public synchronized void produce(Object o) {
         while (buffer!=null) {
6             wait()
         }
         buffer = o
8         notifyAll()
     }

```

¹The idea of using `age` is from [Har98] which uses it in an attempt to propose a fair solution for readers/writers. Unfortunately, the proposed solution is not fair (after an `endWrite` operation, a writer could steal the lock even though there are waiting readers).

```

10     }
12     public synchronized Object consume() {
13         while (buffer==null) {
14             wait()
15         }
16         Object temp = buffer
17         buffer=null
18         notifyAll()
19         return temp
20     }
21 }
22 PC pc = new PC()
23
24 10.times {
25     Thread.start {
26         println (Thread.currentThread().getId()+" consumes")
27         pc.consume()
28     }
29 }
30
31 10.times {
32     Thread.start {
33         println (Thread.currentThread().getId()+" produces")
34         pc.produce((new Random()).nextInt(33))
35     }
36 }

```

Replacing each of the two `notifyAll()` with `notify()` leads to an incorrect solution where one can end up having a producer and consumer both blocked in the wait-set. Hint: C1,C2,P1,P2. This pitfall is called the lost-wakeup problem.

Disadvantages:

- Use multiple condition variables
- Condition variables.

3.3 Readers/Writers

Naive solution. Correct but unfair on writers.

```

import java.util.concurrent.locks.*
2
3 class RW {
4     private int readers;
5     private int writers;
6     static final Lock lock = new ReentrantLock();
7     static final Condition okToRead = lock.newCondition();
8     static final Condition okToWrite = lock.newCondition();
9
10    RW() {
11        readers=0;
12        writers=0;
13    }
14
15    void start_read() {
16        lock.lock();
17        try {

```

```
18     while (writers>0) {
19         okToRead.await();
20     }
21     readers++;
22 } finally {
23     lock.unlock();
24 }
25 }
26
27 void stop_read() {
28     lock.lock();
29     try {
30         readers--;
31         if (readers==0) {
32             okToWrite.signal();
33         }
34     } finally {
35         lock.unlock();
36     }
37 }
38
39 void start_write(Object item) {
40     lock.lock();
41     try {
42         while (readers>0 || writers>0) {
43             okToWrite.await();
44         }
45         writers++;
46     } finally {
47         lock.unlock();
48     }
49 }
50
51 void stop_write() {
52     lock.lock();
53     try {
54         writers--;
55         okToWrite.signal();
56         okToRead.signalAll();
57     } finally {
58         lock.unlock();
59     }
60 }
61
62 }
63
64 RW rw = new RW();
65
66 r = { //R
67     Random r = new Random();
68     rw.start_read();
69     println Thread.currentThread().getId()+" reading..."
70     Thread.sleep(r.nextInt(1000));
71     println Thread.currentThread().getId()+" done reading..."
72
73     rw.stop_read();
74 }
```

```

76 w = { //W
    Random r = new Random();
78    rw.start_write();
    println Thread.currentThread().getId()+" writing..."
80    Thread.sleep(r.nextInt(1000));
    println Thread.currentThread().getId()+" done writing..."
82    rw.stop_write();
    }
84
200.times {
86    Thread.start(r)
    Thread.start(w)
88 }

```

Checking for waiting writers. Places priority on writers but unfair on readers.

```

import java.util.concurrent.locks.*

2
class RW {
4    private int readers;
    private int writers;
6    private int writers_waiting;
    static final Lock lock = new ReentrantLock();
8    static final Condition okToRead = lock.newCondition();
    static final Condition okToWrite = lock.newCondition();

10
    RW() {
12        readers=0;
        writers=0;
14        writers_waiting=0;
    }

16
    void start_read() {
18        lock.lock();
        try {
20            while (writers>0 || writers_waiting>0) {
                okToRead.await();
22            }
            readers++;
24        } finally {
            lock.unlock();
26        }
    }

28
    void stop_read() {
30        lock.lock();
        try {
32            readers--;
            if (readers==0) {
34                okToWrite.signal();
            }
36        } finally {
            lock.unlock();
38        }
    }

40
    void start_write(Object item) {
42        lock.lock();

```



```

44     try {
45         while (readers>0 || writers>0) {
46             writers_waiting++;
47             okToWrite.await();
48             writers_waiting--;
49             writers++;
50         } finally {
51             lock.unlock();
52         }
53     }
54
55     void stop_write() {
56         lock.lock();
57         try {
58             writers--;
59             okToWrite.signal();
60             okToRead.signalAll();
61         } finally {
62             lock.unlock();
63         }
64     }
65 }

```

An attempt at a fair solution to RW is presented in Listing 3.3. One situation that may lead to deadlock is: W1,R1,W2. Another is: R1, W1, R2.

```

1  import java.util.concurrent.locks.*
2
3  class RW {
4      private int readers;
5      private int writers;
6      private int writers_waiting;
7      private int readers_waiting;
8      static final Lock lock = new ReentrantLock();
9      static final Condition okToRead = lock.newCondition();
10     static final Condition okToWrite = lock.newCondition();
11
12     RW() {
13         readers=0;
14         writers=0;
15         writers_waiting=0;
16         readers_waiting=0;
17     }
18
19     void start_read() {
20         lock.lock();
21         try {
22             while (writers>0 || writers_waiting>0) {
23                 readers_waiting++;
24                 okToRead.await();
25                 readers_waiting--;
26             }
27             readers++;
28         } finally {
29             lock.unlock();
30         }
31     }

```

```

33     void stop_read() {
34         lock.lock();
35         try {
36             readers--;
37             if (readers==0) {
38                 okToWrite.signal();
39             }
40         } finally {
41             lock.unlock();
42         }
43     }

44     void start_write(Object item) {
45         lock.lock();
46         try {
47             while (readers>0 || writers>0 || readers_waiting>0) {
48                 writers_waiting++;
49                 okToWrite.await();
50                 writers_waiting--;
51             }
52             writers++;
53         } finally {
54             lock.unlock();
55         }
56     }

57     void stop_write() {
58         lock.lock();
59         try {
60             writers--;
61             okToWrite.signal();
62             okToRead.signalAll();
63         } finally {
64             lock.unlock();
65         }
66     }
67 }
68
69 }

```

Listing 3.3: Incorrect attempt at a fair solution to RW; may deadlock

If we replace `stop_write` with the following code, then our solution may deadlock. Hint: Consider W1,R1,W2.

```

1 void stop_write() {
2     lock.lock();
3     try {
4         writers--;
5         if (readers_waiting==0) {
6             okToWrite.signal();
7         } else {
8             okToRead.signalAll();
9         }
10    } finally {
11        lock.unlock();
12    }
13 }

```

Chapter 4

Promela

4.1 Syntax

We begin with a brief introduction to Promela through a series of examples.

```
1 byte n=0;
3 active proctype P() {
    n=1;
5     printf("P has pid %d. n=%d\n",_pid,n)
    };
7
    active proctype Q() {
9         n=2;
        printf("Q has pid %d. n=%d\n",_pid,n)
11    }
```

Executing Promela code is referred to as a "simulation run of the model".

```
1 $ spin eg.pml
    Q has pid 1. n=2
3     P has pid 0. n=2
2 processes created
```

bash

By default, during simulation runs, SPIN arranges for the output of each active process to appear in a different column: the pid number is used to set the number of tab stops used to indent each new line of output that is produced by a process. You can use the `-t` option to suppress indentation.

```

$ spin -T eg.pml
2 P has pid 0. n=1
  Q has pid 1. n=1
4 2 processes created

```

bash

Semicolon is a separator, not a terminator.

4.1.1 Examples involving Loops

```

byte sum=0;
2
active proctype P() {
4   byte i=0;
   do
6     :: i>10 -> break
     :: else ->
8       sum = sum + i;
       i++;
10  od;
12  printf("The sum of the first 10 numbers is %d\n",sum)
}

```

The following example is one of an infinite loop. Run it and note also how SPIN reports overflows errors.

```

1 byte i=0;
3
active proctype P() {
4   do
5     :: i++;
     printf("Value of i: %d\n. ",i)
7   od
}

```

An example using a for loop:

```

byte sum=0;
2
active proctype P() {
4   byte i;
   for (i:1..10) {
6     sum = sum + i
   }
8
10  printf("The sum of the first 10 numbers is %d\n",sum)
}

```

4.1.2 Expressions as blocking commands

```

byte c=0;
2 finished = 0;
proctype P() {
4   c++;
   finished++
6 }

8 proctype Q() {
   c++;
10  finished++
}

12 init {
14   atomic {
       run P();
16   run Q()
   };
18   finished==2;
   printf("c is %d\n",c)
20 }

```

Equivalently, one may do the following:

```

byte c=0;
2
proctype P() {
4   c++
}

6
proctype Q() {
8   c++
}

10
init {
12   atomic {
       run P();
14   run Q()
   };
16   _nr_pr==1;
   printf("c is %d\n",c)
18 }

```

However, the following variation does not have the expected outcome. When a process terminates, it can only die and make its `_pid` number available for the creation of another process, if and when it has the highest `_pid` number in the system. This means that processes can only die in the reverse order of their creation (in stack order).

```

active proctype P() {
2   printf("A");
}

4
active proctype Q() {
6   printf ("B");
}

8
init {
10  printf("Pr %d",_nr_pr);

```

```

12  _nr_pr==1;
    printf("Done")
}

```

termination.pml

For example, consider what happens if we simulate a run:

```

1  $ spin termination.pml
    A          B          Pr 3          timeout
3  #processes: 3
    3:   proc  2 (:init::1) termination.pml:11 (state 2)
5  3:   proc  1 (Q:1) termination.pml:7 (state 2) <valid end state>
    3:   proc  0 (P:1) termination.pml:3 (state 2) <valid end state>
7  3 processes created

```

bash

It deadlocks at line 11 (`_nr_pr==1`) of the file `termination.pml`. This boolean expression is blocked since processes 0 and 1 cannot terminate until 2 does. If we attempt to verify this program we will obtain an invalid end-state error at line 11.

4.1.3 Macros

Semaphores can be modeled in Promela using an `inline` definition. An inline definition works much like a preprocessor macro, in the sense that it just defines a replacement text for a symbolic name, possibly with parameters.

```

byte s=0;

2  inline acquire(s) {
4  atomic {
    s>0 -> s--
6  }
}

8  inline release(s) {
10 s++
}

12 /* AB after CD */
14 proctype P() {
    acquire(s);
16 printf("A");
    printf("B")
18 }

20 proctype Q() {
    printf("C");
22 printf("D");
    release(s)
24 }

26 init {
    atomic {
28 run P();
    run Q()
}

```

```

30 }
    }

```

Problems if you drop the "atomic" in "acquire":

```

1  byte s=1;
   byte c=0;

3

   inline acquire(s) {
5     s>0 -> s--
   }

7

   inline release(s) {
9     s++
   }

11

   proctype P() {
13     acquire(s);
       c++;
15 }

17 proctype Q() {
       acquire(s);
19     c++;
   }

21 /\ AB after CD *\/ /\
/* proctype P() { */
23 /* acquire(s); */
/* printf("A"); */
25 /* printf("B") */

27 /* } */

29 /* proctype Q() { */

31 /* printf("C"); */
/* printf("D"); */
33 /* release(s) */
/* } */

35

   init {
37     atomic {
       run P();
39     run Q()
   }
41     (_nr_pr==1);
       printf("C is %d ",c)
43 }

```

Exercise: would executing lines 7-8 and 18-19 in atomic block avoid deadlock? What about inverting lines 7 and 8 and then placing them in an atomic block (and likewise with lines 18 and 19)?

```
bool wantP = false;
2 bool wantQ = false;
byte cs=0;
4
proctype P() {
6   do
      :: wantP = true;
8     !wantQ;
      cs++;
10    assert (cs==1);
      cs--;
12    wantP=false
   od
14 }

16 proctype Q() {
   do
18     :: wantQ = true;
      !wantP;
20     cs++;
      assert (cs==1);
22     cs--;
      wantQ=false
24   od
   }
26
28 init {
   atomic {
30     run P();
      run Q()
   }
32 }
```

Figure 4.1: Attempt III in Promela


```

byte ticket=0;
2 byte mutex=1;

4 inline acquire(sem) {
    atomic {
6     sem>0 -> sem--
    }
8 }

10 inline release(sem) {
    sem++
12 }

14 active [5] proctype Jets() {
    acquire(mutex);
16     acquire(ticket);
    acquire(ticket)
18     release(mutex)
    }

20
22 active [5] proctype Patriot() {
    release(ticket);
    }

```

Figure 4.2: Solution to Bar Problem in Promela

4.2 Assertion-Based Model Checking

4.2.1 The Bar Problem Revisited

Listing 4.2 presents the solution to the Bar Problem in Promela. We'll verify that this solution is correct in the sense of upholding the problem invariant, namely that there at least two patriots fans for every jets fan. Before doing so, however, let us first run a simulation of this model.

```

1 > spin bar.pml
    timeout
3 #processes: 5
    ticket = 0
5     mutex = 0
    23: proc 4 (Jets:1) bar.pml:4 (state 4)
    23: proc 3 (Jets:1) bar.pml:4 (state 4)
    23: proc 2 (Jets:1) bar.pml:19 (state 15) <valid end state>
    23: proc 1 (Jets:1) bar.pml:19 (state 15) <valid end state>
    23: proc 0 (Jets:1) bar.pml:4 (state 12)
11 10 processes created

```

bash

The `timeout` indicates that the simulation did not run to completion, it got stuck at a state that is not a valid end state. In other words, it reached a deadlock. From the output above we can see that indeed there are three processes that are deadlocked: 0, 3 and 4. The fact that they are all stuck at line 4 means they are blocked at an `acquire`. Since there are no available permits in `mutex`, clearly processes 3 and 4 are blocked on the `acquire(mutex)` and 0 at the second

acquire(ticket).

A process that terminates must do so after executing its last instruction, otherwise it is said to be in an invalid end state. SPIN checks for this by default. One can insert end state labels to indicate that if execution reaches a certain point and fails to terminate, this should not be considered as an invalid end state. Such valid end state labels must be prefixed with the word `end`. For example, if we replaced the `acquire` operation in 4.2 with the following one:

```

inline acquire(permits) {
2   skip;
end1:
4   atomic {
      permits>0;
6   permits--
      }
8 }

```

then the end states mentioned above are no longer reported as such:

```

> spin bar.pml
2   timeout
#processes: 5
4   ticket = 0
   mutex = 0
6   34:   proc  4 (Jets:1) bar.pml:7 (state 4) <valid end state>
   34:   proc  3 (Jets:1) bar.pml:7 (state 4) <valid end state>
8   34:   proc  2 (Jets:1) bar.pml:22 (state 18) <valid end state>
   34:   proc  1 (Jets:1) bar.pml:7 (state 14) <valid end state>
10  34:   proc  0 (Jets:1) bar.pml:22 (state 18) <valid end state>
10 processes created

```

bash

Let us get back to the task of verifying that the solution is correct. In order to do so we add two counters. Listing 4.2.1 exhibits the updated code.

```

1  byte mutex=1;
   byte ticket=0;
3  byte j=0;
   byte p=0;
5
6  inline acquire(permits) {
7   skip;
end1:
9   atomic {
      permits>0;
11  permits--
      }
13 }
14
15 inline release(permits) {
16   permits++
17 }
18
19 active [5] proctype Jets() {
21   acquire(mutex);

```

```

23   acquire(ticket);
24   acquire(ticket);
25   release(mutex)
26   j++;
27   assert (j*2<=p)
28 }
29
30 active [5] proctype Patriots() {
31   release(ticket)
32   p++;
33   assert (j*2<=p)
34 }

```

We now verify that our solution is correct.

```

1  $ spin -a bar.pml
2  $ gcc -o pan pan.c
3  $ ./pan

5  pan:1: assertion violated ((j*2)<=p) (at depth 34)
6  pan: wrote bar.pml.trail
7
8  (Spin Version 6.5.1 -- 20 December 2019)
9  Warning: Search not completed
10         + Partial Order Reduction
11
12  Full statespace search for:
13      never claim          - (none specified)
14      assertion violations  +
15      acceptance cycles    - (not selected)
16      invalid end states   +
17
18  State-vector 92 byte, depth reached 47, errors: 1
19      18104 states, stored
20      18718 states, matched
21      36822 transitions (= stored+matched)
22      0 atomic steps
23  hash conflicts:          147 (resolved)
24
25  Stats on memory usage (in Megabytes):
26      2.072   equivalent memory usage for states (stored*(State-vector + overhead))
27      1.071   actual memory usage for states (compression: 51.69%)
28              state-vector as stored = 34 byte + 28 byte overhead
29      128.000 memory used for hash table (-w24)
30      0.534   memory used for DFS stack (-m10000)
31      129.511 total actual memory usage
32
33  pan: elapsed time 0.02 seconds
34  pan: rate      905200 states/second

```

bash

It seems that this is not the case since an assertion violation is reported. An inspection of the offending trail shows that when the patriots perform a `release(ticket)` but before incrementing

the `p` counter, a jets fan can go in. There are two ways we can fix our code. One is to increment the `p` counter before performing the release. Another one is to perform the release and increment the counter in one atomic block.

4.2.2 The MEP Problem

Consider the code for Dekker's solution to the MEP from Fig. ???. The Promela code is listed in Fig. ???. We have inserted a variable `cs` to help count when a process enters its critical section. Note how the `await` in line 12 has been coded as a do-loop: we want this loop to cycle while it waits for the condition to hold.

```

1  int turn = 1;
2  boolean wantP = false;
3  boolean wantQ = false;
4
5  Thread.start { //P
6      while (true) {
7          // non-CS
8          wantP = true
9          while wantQ
10             if (turn == 2) {
11                 wantP = false
12                 await (turn==1)
13                 wantP = true
14             }
15             // CS
16             turn = 2
17             wantP = false
18             // non-CS
19         }
20     }
21
22     Thread.start { //Q
23         while (true) {
24             // non-CS
25             wantQ = true
26             while wantP
27                 if (turn == 1) {
28                     wantQ = false
29                     await (turn==2)
30                     wantQ = true
31                 }
32                 // CS
33                 turn = 1
34                 wantQ = false
35                 // non-CS
36             }
37         }
38     }
39 }

```

```

1  bool wantp = false;
2  bool wantq = false;
3  byte turn = 1;
4  byte cs=0;
5
6  active proctype P() {
7      do

```

```

9      :: wantp = true;
      do
10         :: !wantq -> break;
11         :: else ->
            if
12             :: (turn == 2) ->
                wantp = false;
13             do
14                 :: turn==1 -> break
15                 :: else
16                 od;
17                 wantp = true
18             :: else /* leaves if, if turn<>2 */
19             fi
20         od;
21         cs++;
22         assert(cs==1);
23         cs--;
24         wantp = false;
25         turn = 2
26     od
27 }

31 active proctype Q() {
32     do
33         :: wantq = true;
34         do
35             :: !wantp -> break;
36             :: else ->
37                 if
38                     :: (turn == 1) ->
39                     wantq = false;
40                     do
41                         :: turn==2 -> break
42                         :: else
43                         od;
44                     wantq = true
45                 :: else /* leaves if, if turn<>2 */
46                 fi
47             od;
48             cs++;
49             assert(cs==1);
50             cs--;
51             wantq = false;
52             turn = 1
53         od
54 }

```

```

$ spin -a dekker.pml
2 $ gcc -o pan pan.c
$ ./pan

4
(Spin Version 6.5.1 -- 20 December 2019)
6 + Partial Order Reduction

8 Full statespace search for:

```

```

10      never claim          - (none specified)
      assertion violations  +
12      acceptance cycles   - (not selected)
      invalid end states   +

14 State-vector 28 byte, depth reached 74, errors: 0
      172 states, stored
16      173 states, matched
      345 transitions (= stored+matched)
18      0 atomic steps
hash conflicts:          0 (resolved)

20 Stats on memory usage (in Megabytes):
22      0.009   equivalent memory usage for states (stored*(State-vector + overhead))
      0.287   actual memory usage for states
24      128.000 memory used for hash table (-w24)
      0.534   memory used for DFS stack (-m10000)
26      128.730 total actual memory usage

28 unreachable in proctype P
30      dekker.pml:29, state 28, "-end-"
      (1 of 28 states)
32 unreachable in proctype Q
34      dekker.pml:54, state 28, "-end-"
      (1 of 28 states)

36 pan: elapsed time 0 seconds

```

bash

4.3 Non-Progress Cycles

SPIN can check for some simple liveness properties without the need to use Temporal Logic. An infinite computation that does not include infinitely many occurrences of a progress state is called a non-progress cycle. We illustrate this feature by showing that Dekker's algorithm enjoys absence of livelock.

Consider

```

byte x=1;

2 active proctype P() {
4
      do
6          :: x==1 -> x=2;
          :: x==2 -> x=1;
8      od
}

```

Consider

```

1 byte x=1;

3 active proctype P() {

5     do
6         :: x==1 -> x=2;
7         :: x==2 -> progress1: x=1;
8     od
9 }

```

Consider

```

1 byte x=1;

3 active proctype P() {

5     do
6         :: x==1 -> x=2;
7         :: x==2 -> progress1: x=1;
8         :: x==2 -> x=1;
9     od
10 }

```

We would like to verify that this attempt at solving the MEP problem does not enjoy absence of livelock. For that we insert progress labels just before entering the CS.

```

bool wantP=false;
bool wantQ=false;

4 proctype P() {
5     do
6         :: wantP=true;
7         do
8             :: wantQ==false -> break
9             :: else
10                od;
11 progress1:
12         wantP=false
13     od
14 }

16 proctype Q() {
17     do
18         :: wantQ=true;
19         do
20             :: wantP==false -> break
21             :: else
22                od;
23 progress2:
24         wantQ=false
25     od
26 }

28 init {
29     atomic {
30         run P();
31         run Q()
32     }
33 }

```

```
}

```

Selecting Non-Progress in the drop down list and then verifying, SPIN reports a non-progress cycle:

```

1 2 Q:1 1) wantQ = 1
   Process Statement      wantQ
3 1 P:1 1) wantP = 1      1
   Process Statement      wantP      wantQ
5 2 Q:1 1) else           1          1
   <<<<<START OF CYCLE>>>>>
7 2 Q:1 1) else           1          1
   1 P:1 1) else           1          1
9 2 Q:1 1) else           1          1
spin: trail ends after 15 steps

```

spin

```

bool wantp = false;
2 bool wantq = false;
byte turn = 1;
4
active proctype P() {
6   do
8     :: wantp = true;
10    do
12      :: !wantq -> break;
14      :: else ->
16        if
18          :: (turn == 2) ->
20            wantp = false;
22            do
24              :: turn==1 -> break
26              :: else
28                od;
30                wantp = true
32              :: else /* leaves if, if turn<>2 */
34                fi
36            od;
38  progressP:
39    wantp = false;
40    turn = 2
41  od
42 }
43
44 active proctype Q() {
45   do
46     :: wantq = true;
47     do
48       :: !wantp -> break;
49       :: else ->
50         if
51           :: (turn == 1) ->
52             wantq = false;
53             do
54               :: turn==2 -> break
55               :: else

```



```

40         od;
         wantq = true
42     :: else /* leaves if, if turn<>2 */
         fi
44     od;
progressQ:
46     wantq = false;
         turn = 1
48     od
}

```

“Weak Fairness” should be enabled. Weak fairness means that each statement that becomes enabled and remains enabled thereafter will eventually be scheduled. Consider the example below [?]:

```

1 byte x=0;

3 active proctype P() {
    do
5         :: true -> x = 1 - x;
    od
7 }

9 active proctype Q() {
    do
11        :: true -> progress1: x = 1 - x;
    od
13 }

```

It is possible that Q makes no progress if Q is never scheduled for execution. Weak fairness guarantees that it eventually will. Verify this in SPIN by first enabling weak fairness and then disabling it. In the former case no errors are reported, but in the latter a non-progress cycle is reported:

```

1 0 P:1 1) 1
<<<<START OF CYCLE>>>>
3 0 P:1 1) x = (1-x)
Process Statement      x
5 0 P:1 1) 1           1
0 P:1 1) x = (1-x)     1
7 0 P:1 1) 1           0

```

spin

Consider the code for Attempt IV

```

1 bool wantP = false, wantQ = false;

3 active proctype P() {
    do
5         :: wantP = true;
         do
7             :: wantQ -> wantP = false; wantP = true
             :: else -> break
         od;
9         wantP = false

```

```

11     od
12 }
13
14 active proctype Q() {
15     do
16         :: wantQ = true;
17         do
18             :: wantP -> wantQ = false; wantQ = true
19             :: else -> break
20         od;
21     wantQ = false
22     od
23 }

```

We know that it does not enjoy freedom from starvation. Freedom from starvation would mean that both P and Q enter their CS infinitely often. We can verify that it does not enjoy freedom from starvation by inserting a progress label in the critical section of P, selecting Non-Progress in the drop down list and then verifying.

```

1     bool wantP = false, wantQ = false;
2
3     active proctype P() {
4         do
5             :: wantP = true;
6             do
7                 :: wantQ -> wantP = false; wantP = true
8                 :: else -> break
9             od;
10        progress1:
11            wantP = false
12            od
13    }
14
15    active proctype Q() {
16        do
17            :: wantQ = true;
18            do
19                :: wantP -> wantQ = false; wantQ = true
20                :: else -> break
21            od;
22        progress2:
23            wantQ = false
24            od
25    }

```

Here is the output from SPIN

```

1  1 Q:1  1)  wantQ = 1
   Process Statement      wantQ
3  1 Q:1  1)  else          1
   1 Q:1  1)  wantQ = 0      1
5  0 P:1  1)  wantP = 1      0
   Process Statement      wantP    wantQ
7  1 Q:1  1)  wantQ = 1      1        0
   1 Q:1  1)  wantP          1        1
9  0 P:1  1)  wantQ          1        1
   <<<<START OF CYCLE>>>>
11 1 Q:1  1)  wantQ = 0      1        1

```

```

1 Q:1 1) wantQ = 1 1 0
13 1 Q:1 1) wantP 1 1
0 P:1 1) wantP = 0 1 1
15 1 Q:1 1) wantQ = 0 0 1
1 Q:1 1) wantQ = 1 0 0
17 1 Q:1 1) else 0 1
0 P:1 1) wantP = 1 0 1
19 0 P:1 1) wantQ 1 1
1 Q:1 1) wantQ = 0 1 1
21 0 P:1 1) wantP = 0 1 0
1 Q:1 1) wantQ = 1 0 0
23 1 Q:1 1) else 0 1
1 Q:1 1) wantQ = 0 0 1
25 0 P:1 1) wantP = 1 0 0
1 Q:1 1) wantQ = 1 1 0
27 1 Q:1 1) wantP 1 1
Process Statement wantP wantQ
29 0 P:1 1) wantQ 1 1
spin: trail ends after 50 steps

```

4.3.1 The Zoo Problem Revisited

Consider the Zoo Problem discussed in Exercise ??:

In a zoo there is a common feeding area for exotic mice and exotic felines. The feeding area has a common feeding lot that holds up to 2 animals. The feeding area can be used by both mice and felines, but cannot be used by mice and felines at the same time for obvious reasons.

A solution in Promela is given in Listing 4.3.

Exercise 4.3.1. Show that if lines are removed, then deadlock is possible. Explain the deadlock situation that can arise.

Exercise 4.3.2. Show, using assertions, that there cannot be felines feeding, if there are mice feeding and, likewise, there cannot be mice feeding, if there are felines feeding.

Exercise 4.3.3. You are asked to show that there can be at most two mice and at most two felines in the feeding lot. Since you already know that both cannot be feeding at the same time, you propose the following assertions (only the portion of code that is modified is shown) in the Mouse process:

```

acquire(feedinglot);
2 // access feeding lot
assert (mice<3);
4 release(feedinglot);

```

Similarly for the Feline process:

```

acquire(feedinglot);
2 // use feeding lot
assert (felines<3);
4 release(feedinglot);

```

```

byte mice = 0;
2 byte felines = 0;
byte mutexMice = 1;
4 byte mutexFelines = 1;
byte feedinglot = 2;
6 byte mutex = 1;

8 inline acquire(sem) {
    atomic {
10     sem>0;
        sem--
12     }
    }

14 inline release(sem) {
16     sem++
    }

18 active [3] proctype Mouse() {
20     acquire(mutex);
    acquire(mutexMice);
22     mice++;
    if
24     :: mice==1 -> acquire(mutexFelines);
        :: else -> skip;
26     fi
    release(mutexMice);
28     release(mutex);

30     acquire(feedinglot);
    // access feeding lot
32     release(feedinglot);

34     acquire(mutexMice);
    mice--;
36     if
        :: mice==0 -> release(mutexFelines);
        :: else -> skip;
38     fi
    release(mutexMice);
40 }

42 active [3] proctype Feline() {
44     acquire(mutex);
    acquire(mutexFelines);
46     felines++;
    if
48     :: felines==1 -> acquire(mutexMice);
        :: else -> skip;
50     fi
    release(mutexFelines);
52     release(mutex);

54     acquire(feedinglot);
    // use feeding lot
56     release(feedinglot);

58     acquire(mutexFelines);
    felines--;
60     if
v0.015: felines==0 -> release(mutexMice);
62     :: else -> skip;
        fi
64     release(mutexFelines);
66 }

```

CP Notes

Figure 4.3: Zoo Problem in Promela

Unfortunately, when you verify this with Spin, it reports an assertion violation.

- 1. Explain why.*
- 2. Replace the assertions with new ones that address the problem correctly.*

Chapter 5

Solution to Selected Exercises

Section ??

Answer 5.0.1 (Exercise ??). *jj*

Section ??

Bibliography

- [Car96] Tom Cargill. Specific notification for java thread synchronization. www.dre.vanderbilt.edu/%7Eeschmidt/PDF/specific-notification.pdf, 1996.
- [Har98] Stephen Hartley. Concurrent Programming: The Java Programming Language. Oxford University Press, 1998.