

CS579: Foundations of Cryptography

Spring 2023

Symmetric-Key Encryption

Instructor: **Nikos Triandopoulos**



Computational security

The big picture

- ◆ we formally defined and constructed a perfectly secure cipher
- ◆ this encryption has some drawbacks
 - ◆ e.g., it employs a very large key
- ◆ by Shannon's Theorem, such limitations are unavoidable



Now, what?

Our approach: Relax “perfectness”

Initial model / abstraction

- ◆ **perfect secrecy** / security requires that
 - ◆ absolutely **no information is leaked** about the plaintext
 - ◆ to adversaries that **unlimited computational power**

Refined model / abstraction

- ◆ consider a relaxed notion of security, called **computational security**, where an encryption scheme is (for all **practical** purposes) secure even if
 - ◆ **a tiny amount of information is leaked** about the plaintext (e.g., w/ prob. 2^{-60})
 - ◆ to adversaries with **bounded computational power** (e.g., attacker invests 200ys)

Computational security

- ◆ to be contrasted against information-theoretic security
- ◆ de facto way in which security is modeled in most cryptographic settings
- ◆ an integral part of modern cryptography w/ rigorous mathematical proofs
- ◆ entails two relaxations
 - ◆ **security is guaranteed against efficient adversaries**
 - ◆ if an attacker invests in sufficiently large resources, it may break the scheme's security
 - ◆ goal: make required resources larger than those available to any realistic attacker!
 - ◆ **adversaries can potentially succeed** (i.e., security guarantees are probabilistic)
 - ◆ with some small probability the scheme is breakable
 - ◆ goal: make success probability sufficiently small so that it can be practically ignored!

Towards a rigorous definition of computational security

- ◆ **Concrete** approach

- ◆ bounds the maximum success probability of any (randomized) adversary running for some specified amount of time or investing a specified amount of resources
- ◆ general security result: *“A scheme is (t, ϵ) -secure if any adversary \mathcal{A} , running for time at most t , succeeds in breaking the scheme with probability at most ϵ ”*
- ◆ need to
 - ◆ define what it means for an adversary to “break” a scheme
 - ◆ specify precisely the resources (e.g., time in seconds using a particular computer, or CPU cycles in a particular available supercomputer architecture)

Examples

- ◆ almost optimal security guarantees
 - ◆ key length n , key space size $|K| = 2^n$
 - ◆ \mathcal{A} running for time t (e.g., CPU cycles) succeeds w/ prob. at most $ct/2^n$
 - ◆ this corresponds to a brute-force type of attack (w/out preprocessing)
- ◆ parameter c models advanced computing methods
 - ◆ c is typically larger than 1: e.g., parallelism can be used, etc.
- ◆ if $c = 1$, $n = 60$, security is enough for attackers running a desktop computer
 - ◆ 4 GHz (4×10^9 cycles/sec), 2^{60} CPU cycles require about 9 years
 - ◆ however, “fastest” available computer runs w/ 2×10^{16} cycles/sec, i.e., in ~ 1 min!
 - ◆ choosing $n=80$ is better: the supercomputer would still need ~ 2 years

Today's recommendations

- ◆ recommended security parameter is $n = 128$
- ◆ large difference between 2^{80} and 2^{128}
 - ◆ #seconds since Big Bang is $\sim 2^{58}$
- ◆ if probability of success (within 1 year of computation) is $1/2^{60}$
 - ◆ it is more likely that Alice and Bob are hit by lightning (and thus don't care much about the breached confidentiality)
 - ◆ an event happening once in 100 years corresponds to probability 2^{-30} of happening at a particular second
- ◆ limitations of the concrete approach
 - ◆ harder to achieve, careful interpretation is needed, what if \mathcal{A} runs for $2t$ or $t/2$?

An alternative (less quantitative) approach

- ◆ **Asymptotic** approach
 - ◆ again, a security parameter n is used (e.g., key length)
 - ◆ **efficient (or realistic or feasible) adversaries** are equated with probabilistic poly-time (PPT) algorithms that run for time that is a polynomial of n
 - ◆ **small probability of success** is equated with success probabilities that are asymptotically smaller than any inverse polynomial in n
 - ◆ general security result: *“A scheme is secure if any PPT adversary \mathcal{A} succeeds in breaking the scheme with at most negligible probability”*

Negligible functions (to capture tiny likelihood)

Typically, they measure the probability of success (of an attacker)

- ◆ Intuitively: very small probability
 - ◆ negligible, can be ignored, it's more likely to be hit by asteroid...
 - ◆ approaches 0 faster than the inverse of any polynomial
 - ◆ notation: **negl**
- ◆ Formally
 - ◆ A function $\mu : \mathbf{N} \rightarrow \mathbf{R}^+$ is negligible if for every positive integer c there exists an integer N such that for all $n > N$, it holds that

$$\mu(n) < 1 / n^c$$

Security parameter

- ◆ Asymptotic approach
 - ◆ general result: “A scheme is **secure** if any **PPT adversary** \mathcal{A} succeeds in breaking the scheme with at most **negligible** probability”
 - ◆ the terms “**negligible**” and “**polynomial**” make sense only if any algorithm (and the adversary \mathcal{A}) takes an additional input 1^n
 - ◆ this is called the **security parameter**
 - ◆ i.e., we consider an infinite sequence of schemes Π parameterized by n
 $\Pi(1), \Pi(2), \dots$
- ◆ e.g., security parameter n equals the key length (i.e., $\{0,1\}^n \rightarrow k$)
 - ◆ \mathcal{A} can always guess k with probability 2^{-n} – a negligible function of n
 - ◆ \mathcal{A} can also enumerate all possible keys k in time 2^n – an exponential time in n

Asymptotic approach: Pros & cons

- ◆ Pros

- ◆ all types of Turing Machines are “equivalent” up to a “polynomial reduction”
- ◆ no need to specify the details of the computational model
- ◆ in analyzing a scheme, the involved formulas get much simpler

- ◆ Cons

- ◆ asymptotic results don't tell us anything about security of the concrete systems

- ◆ In practice

- ◆ we prove formally an asymptotic result; and then
- ◆ argue informally that “the constants are reasonable”
(or calculate them if it is needed)

Negligible functions

Recall: Defining security relaxations

- ◆ Concrete approach
 - ◆ general result: “A scheme is **(t, ϵ) -secure** if any adversary \mathcal{A} , running for **time at most t** , succeeds in breaking the scheme with probability **at most ϵ** ”
- ◆ Asymptotic approach
 - ◆ general result: “A scheme is **secure** if any **PPT adversary** \mathcal{A} succeeds in breaking the scheme with at most **negligible** probability”
 - ◆ PPT algorithm: probabilistic algorithm that runs in time $O(n^c)$ for some c
 - ◆ notation: **poly**

Negligible functions

Typically, they measure the probability of success (of an attacker)

- ◆ Intuitively: very small probability
 - ◆ negligible, can be ignored, it's more likely to be hit by asteroid...
 - ◆ approaches 0 faster than the inverse of any polynomial
 - ◆ notation: **negl**
- ◆ Formally
 - ◆ A function $\mu : \mathbf{N} \rightarrow \mathbf{R}^+$ is negligible if for every positive integer c there exists an integer N such that for all $n > N$, it holds that

$$\mu(n) < 1 / n^c$$

Example of negligible functions

$f(n) =$

- ◆ $1 / n^2 \rightarrow \text{No}$
- ◆ $2^{-n} \rightarrow \text{Yes}$
 - ◆ E.g., for $n > 23$, $f(n) < n^{-5}$
- ◆ $2^{-\sqrt{n}} \rightarrow \text{Yes}$
 - ◆ E.g., for $n > 3500$, $f(n) < n^{-5}$
- ◆ $n^{-\log n} \rightarrow \text{Yes}$
 - ◆ E.g., for $n > 33$, $f(n) < n^{-5}$
- ◆ $1 / n^{10000} \rightarrow \text{No}$

Properties of **poly** and **negl** functions

- ◆ A sum of two polynomials is a polynomial

$$\text{poly} + \text{poly} = \text{poly}$$

- ◆ A product of two polynomials is a polynomial:

$$\text{poly} * \text{poly} = \text{poly}$$

- ◆ A sum of two negligible functions is a negligible function:

$$\text{negl} + \text{negl} = \text{negl}$$

Moreover:

- ◆ A negligible function multiplied by a polynomial is negligible

$$\text{negl} * \text{poly} = \text{negl}$$

Computational secrecy

Recall: Approach in modern cryptography

Formal treatment


- ◆ **fundamental notions** underlying the **design & evaluation** of crypto primitives

Systematic process

- ◆ (A) **formal definitions** (what it means for a crypto primitive to be “secure”?)
- ◆ (B) **precise assumptions** (which forms of attacks are allowed – and which aren’t?)
- ◆ (C) **provable security** (why a candidate instantiation is indeed secure – or not?)

Example: Precise assumptions (1)

- ◆ **adversary**


- ◆ type of attacks – a.k.a. **threat model**  **eavesdropping**
- ◆ **capabilities** (e.g., a priori knowledge, access to information, party corruptions)
- ◆ **limitations** (e.g., bounded memory, passive Vs. active)



Eve may know the a priori distribution of messages sent by Alice



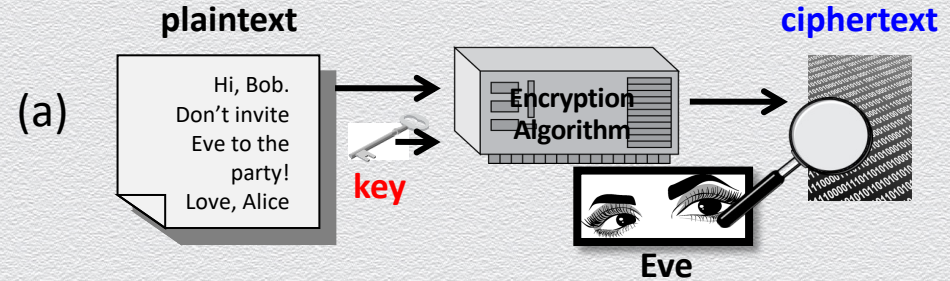
Eve doesn't know/learn the secret k (shared by Alice and Bob)



Example: Possible eavesdropping attacks (1.I)

An attacker may possess a


- ◆ (a) collection of ciphertexts
 - ◆ ciphertext only attack
 - ◆ this will be the **default attack type** when we will next define the concept of perfect security



Example: Possible eavesdropping attacks (1.II)

An attacker may possess a

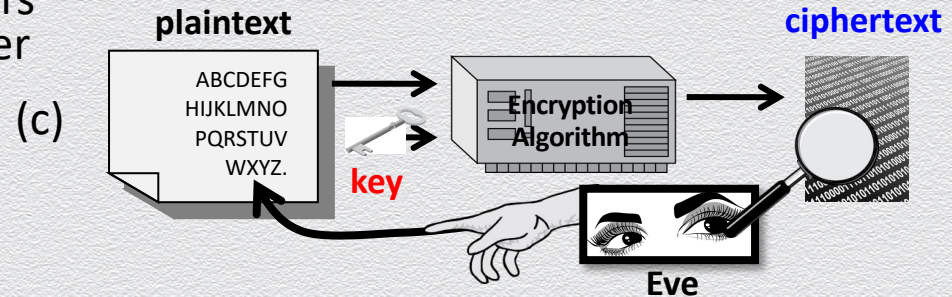
- ◆ (a) collection of ciphertexts
 - ◆ ciphertext only attack
- ◆ (b) collection of plaintext/ciphertext pairs
 - ◆ known plaintext attack



Example: Possible eavesdropping attacks (1.III)

An attacker may possess a


- ◆ (a) collection of ciphertexts
 - ◆ ciphertext only attack
- ◆ (b) collection of plaintext/ciphertext pairs
 - ◆ known plaintext attack
- ◆ (c) collection of plaintext/ciphertext pairs for plaintexts selected by the attacker
 - ◆ chosen plaintext attack



Example: Possible eavesdropping attacks (1.IV)

An attacker may possess a

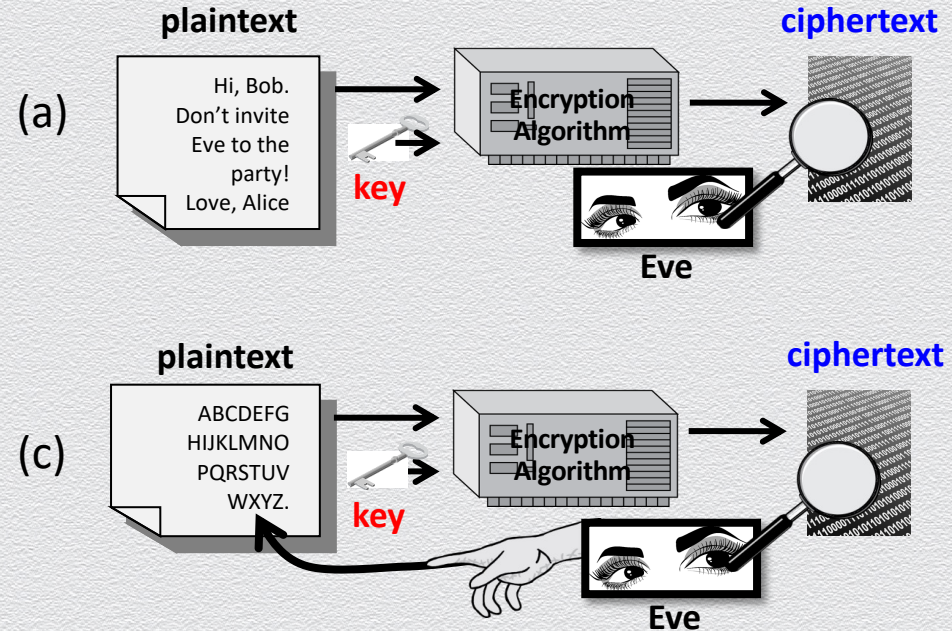
- ◆ (a) collection of ciphertexts
 - ◆ ciphertext only attack
- ◆ (b) collection of plaintext/ciphertext pairs
 - ◆ known plaintext attack
- ◆ (c) collection of plaintext/ciphertext pairs for plaintexts selected by the attacker
 - ◆ chosen plaintext attack
- ◆ (d) collection of plaintext/ciphertext pairs for (plaintexts and) ciphertexts selected by the attacker
 - ◆ chosen ciphertext attack



Recall: Possible eavesdropping attacks

An attacker may possess a

- ◆ (a) collection of ciphertexts
 - ◆ ciphertext only attack
 - ◆ **EAV-attack**
- ◆ (c) collection of plaintext/ciphertext pairs for plaintexts selected by the attacker
 - ◆ chosen plaintext attack
 - ◆ CPA-attack



3 equivalent definitions of perfect EAV-security

1) a posteriori = a priori

For every $\mathcal{D}_{\mathcal{M}}$, $m \in \mathcal{M}$ and $c \in \mathcal{C}$, for which $\Pr[C = c] > 0$, it holds that

$$\Pr[M = m \mid C = c] = \Pr[M = m]$$

2) \mathcal{C} is independent of \mathcal{M}


For every $m, m' \in \mathcal{M}$ and $c \in \mathcal{C}$, it holds that

$$\Pr[\text{Enc}_K(m) = c] = \Pr[\text{Enc}_K(m') = c]$$

3) indistinguishability

For every \mathcal{A} , it holds that

$$\Pr[b' = b] = 1/2$$



From perfect to computational EAV-security

- ◆ **perfect** security: M , $\text{Enc}_K(M)$ are independent
 - ◆ absolutely **no information is leaked** about the plaintext
 - ◆ to adversaries that **unlimited computational power**
- ◆ **computational** security: for all **practical** purposes, M , $\text{Enc}_K(M)$ are independent
 - ◆ **a tiny amount of information is leaked** about the plaintext (e.g., w/ prob. 2^{-60})
 - ◆ to adversaries with **bounded computational power** (e.g., attacker invests 200ys)
- ◆ attacker's **best strategy** remains **ineffective**
 - ◆ **random guess** on secret key; or
 - ◆ **exhaustive search** over key space (**brute force attack**)

A formal, mathematic view of symmetric encryption

A symmetric-key encryption scheme is defined by

- ◆ a **message space** \mathcal{M} , $|\mathcal{M}| > 1$, and a triple (**Gen**, **Enc**, **Dec**)
- ◆ **Gen**: probabilistic key-generation algorithm, defines **key space** \mathcal{K}
 - ◆ $\text{Gen}(1^n) \rightarrow k \in \mathcal{K}$ (security parameter n)
- ◆ **Enc**: probabilistic encryption algorithm, defines **ciphertext space** \mathcal{C}
 - ◆ $\text{Enc}: \mathcal{K} \times \mathcal{M} \rightarrow \mathcal{C}$, $\text{Enc}(k, m) = \text{Enc}_k(m) \rightarrow c \in \mathcal{C}$
- ◆ **Dec**: deterministic encryption algorithm
 - ◆ $\text{Dec}: \mathcal{K} \times \mathcal{C} \rightarrow \mathcal{M}$, $\text{Dec}(k, c) = \text{Dec}_k(c) := m \in \mathcal{M}$ or \perp

Relaxing indistinguishability

Relax the definition of perfect secrecy – that is based on indistinguishability

- ◆ require that $\mathbf{m}_0, \mathbf{m}_1$ are chosen by a **PPT adversary**
- ◆ require that no **PPT adversary** can distinguish $\mathbf{Enc}_k(\mathbf{m}_0)$ from $\mathbf{Enc}_k(\mathbf{m}_1)$

non-negligibly better than guessing


3) indistinguishability

For every \mathcal{A} , it holds that

$$\Pr[\mathbf{b}' = \mathbf{b}] = 1/2 + \text{negl}$$

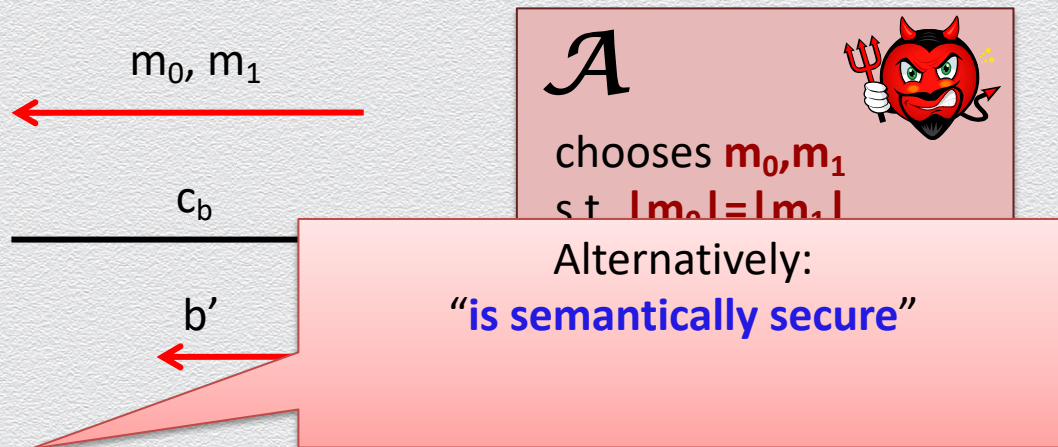
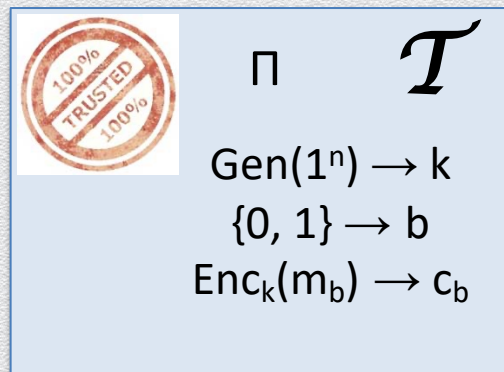
PPT

negl



Game-based definition of computational EAV-security

encryption scheme $\Pi = \{\mathcal{M}, (\text{Gen}, \text{Enc}, \text{Dec})\}$

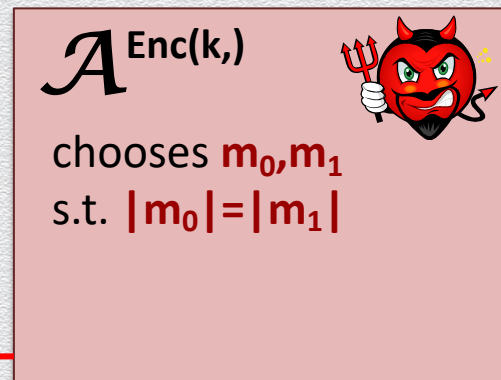

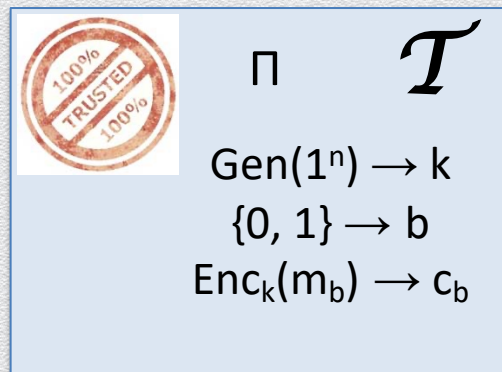


We say that (Enc, Dec) is **EAV-secure** if any PPT adversary \mathcal{A} guesses b correctly with probability at most $0.5 + \epsilon(n)$, where ϵ is a negligible function

i.e., no PPT \mathcal{A} computes b correctly non-negligibly better than randomly guessing

Similarly: CPA-security

encryption scheme $\Pi = \{\mathcal{M}, (\text{Gen}, \text{Enc}, \text{Dec})\}$



We say that (Enc, Dec) is **CPA-secure** if any PPT adversary \mathcal{A} guesses b correctly with probability at most $0.5 + \epsilon(n)$, where ϵ is a negligible function

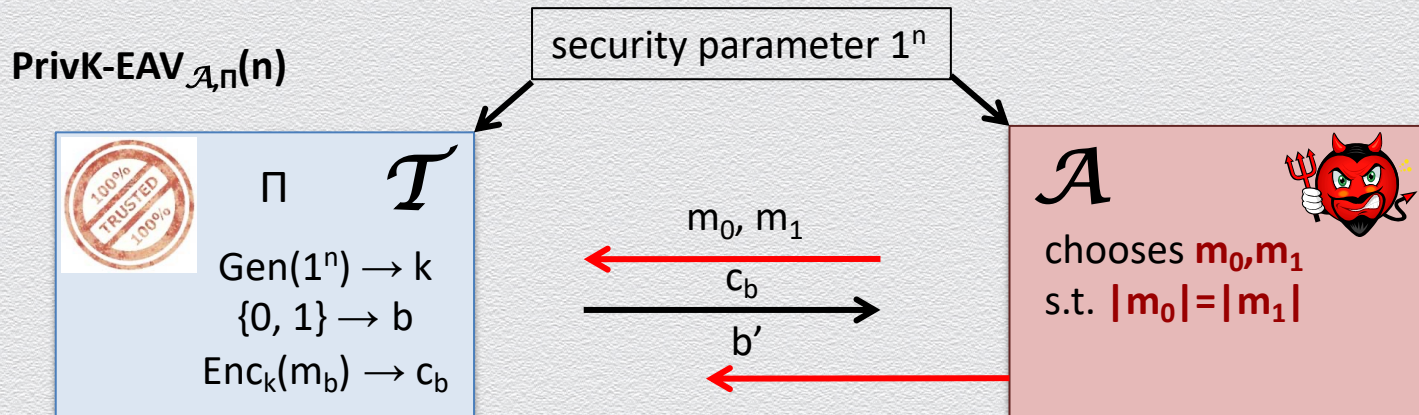
I.e., no PPT \mathcal{A} computes b correctly non-negligibly better than randomly guessing, **even when it learns the encryptions of messages of its choice**

On CPA security

Facts

- ◆ Any encryption scheme that is CPA-secure is also CPA-secure for multiple encryptions
- ◆ **CPA security implies probabilistic encryption – can you see why?**
- ◆ EAV-security for multiple messages implies probabilistic encryption

Recall: EAV-security for secrecy



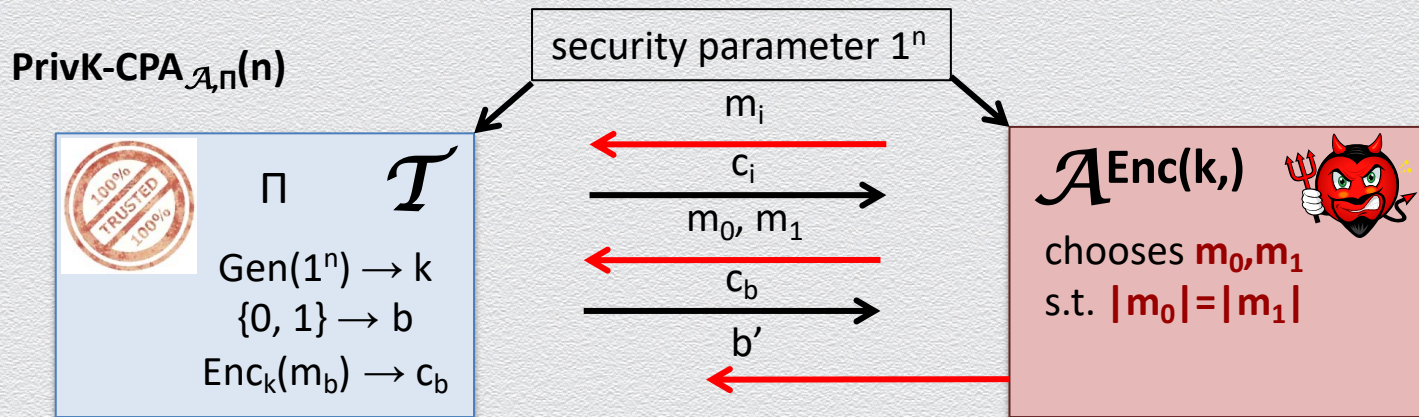
Security definition

We say that $\Pi = (\text{Enc}, \text{Dec})$ is **EAV-secure**, if \forall PPT adversaries \mathcal{A} , \exists a negligible function $\epsilon(n)$ such that it holds that

$$\Pr[\text{PrivK-EAV}_{\mathcal{A},n}(n)=1] \leq \frac{1}{2} + \epsilon(n)$$

I.e., no PPT \mathcal{A} computes b correctly non-negligibly better than randomly guessing

Recall: CPA-security for secrecy




Security definition

We say that $\Pi = (\text{Enc}, \text{Dec})$ is **CPA-secure**, if \forall PPT adversaries \mathcal{A} , \exists a negligible function $\epsilon(n)$ such that it holds that

$$\Pr[\text{PrivK-CPA}_{\mathcal{A},\Pi}(n)=1] \leq \frac{1}{2} + \epsilon(n)$$

I.e., no PPT \mathcal{A} computes b correctly non-negligibly better than randomly guessing,
even when it learns the encryptions of messages of its choice

CPA-security for secrecy for multiple messages



Security definition

We say that $\Pi = (\text{Enc}, \text{Dec})$ is **CPA-secure for multiple encryptions**, if \forall PPT adversaries \mathcal{A} , \exists a negligible function $\varepsilon(n)$ such that it holds that

$$\Pr[\text{PrivK-MCPA}_{\mathcal{A}, \Pi}(n) = 1] \leq \frac{1}{2} + \varepsilon(n)$$

I.e., no PPT \mathcal{A} computes **b** for **many challenge ciphertexts** correctly non-negligibly better than randomly guessing, even when it learns the encryptions of messages of its choice